



🏠 → [Binary data, files](#)

📅 10th October 2019

# ArrayBuffer, binary arrays

In web-development we meet binary data mostly while dealing with files (create, upload, download). Another typical use case is image processing.

That's all possible in JavaScript, and binary operations are high-performant.

Although, there's a bit of confusion, because there are many classes. To name a few:

- `ArrayBuffer` , `Uint8Array` , `DataView` , `Blob` , `File` , etc.

Binary data in JavaScript is implemented in a non-standard way, compared to other languages. But when we sort things out, everything becomes fairly simple.

**The basic binary object is `ArrayBuffer` – a reference to a fixed-length contiguous memory area.**

We create it like this:

```
1 let buffer = new ArrayBuffer(16); // create a buffer of length 16
2 alert(buffer.byteLength); // 16
```



This allocates a contiguous memory area of 16 bytes and pre-fills it with zeroes.

## ⚠️ **ArrayBuffer is not an array of something**

Let's eliminate a possible source of confusion. `ArrayBuffer` has nothing in common with `Array` :

- It has a fixed length, we can't increase or decrease it.
- It takes exactly that much space in the memory.
- To access individual bytes, another "view" object is needed, not `buffer[index]` .

`ArrayBuffer` is a memory area. What's stored in it? It has no clue. Just a raw sequence of bytes.

**To manipulate an `ArrayBuffer` , we need to use a "view" object.**

A view object does not store anything on it's own. It's the "eyeglasses" that give an interpretation of the bytes stored in the `ArrayBuffer` .

For instance:

- **`Uint8Array`** – treats each byte in `ArrayBuffer` as a separate number, with possible values are from 0 to 255 (a byte is 8-bit, so it can hold only that much). Such value is called a "8-bit unsigned integer".
- **`Uint16Array`** – treats every 2 bytes as an integer, with possible values from 0 to 65535. That's called a "16-bit unsigned integer".

- **Uint32Array** – treats every 4 bytes as an integer, with possible values from 0 to 4294967295. That's called a "32-bit unsigned integer".
- **Float64Array** – treats every 8 bytes as a floating point number with possible values from  $5.0 \times 10^{-324}$  to  $1.8 \times 10^{308}$ .

So, the binary data in an `ArrayBuffer` of 16 bytes can be interpreted as 16 "tiny numbers", or 8 bigger numbers (2 bytes each), or 4 even bigger (4 bytes each), or 2 floating-point values with high precision (8 bytes each).

new ArrayBuffer(16)																
Uint8Array	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15
Uint16Array	0		1		2		3		4		5		6		7	
Uint32Array	0				1				2				3			
Float64Array	0								1							

`ArrayBuffer` is the core object, the root of everything, the raw binary data.

But if we're going to write into it, or iterate over it, basically for almost any operation – we must use a view, e.g:

```

1 let buffer = new ArrayBuffer(16); // create a buffer of length 16
2
3 let view = new Uint32Array(buffer); // treat buffer as a sequence of 32-bit integers
4
5 alert(Uint32Array.BYTES_PER_ELEMENT); // 4 bytes per integer
6
7 alert(view.length); // 4, it stores that many integers
8 alert(view.byteLength); // 16, the size in bytes
9
10 // let's write a value
11 view[0] = 123456;
12
13 // iterate over values
14 for(let num of view) {
15   alert(num); // 123456, then 0, 0, 0 (4 values total)
16 }
```

## TypedArray

The common term for all these views (`Uint8Array`, `Uint32Array`, etc) is `TypedArray`. They share the same set of methods and properties.

They are much more like regular arrays: have indexes and iterable.

A typed array constructor (be it `Int8Array` or `Float64Array` , doesn't matter) behaves differently depending on argument types.

There are 5 variants of arguments:

```
1 new TypedArray(buffer, [byteOffset], [length]);
2 new TypedArray(object);
3 new TypedArray(typedArray);
4 new TypedArray(length);
5 new TypedArray();
```

1. If an `ArrayBuffer` argument is supplied, the view is created over it. We used that syntax already.

Optionally we can provide `byteOffset` to start from (0 by default) and the `length` (till the end of the buffer by default), then the view will cover only a part of the `buffer` .

2. If an `Array` , or any array-like object is given, it creates a typed array of the same length and copies the content.

We can use it to pre-fill the array with the data:

```
1 let arr = new Uint8Array([0, 1, 2, 3]);
2 alert( arr.length ); // 4, created binary array of the same length
3 alert( arr[1] ); // 1, filled with 4 bytes (unsigned 8-bit integers) with 0
```

3. If another `TypedArray` is supplied, it does the same: creates a typed array of the same length and copies values. Values are converted to the new type in the process, if needed.

```
1 let arr16 = new Uint16Array([1, 1000]);
2 let arr8 = new Uint8Array(arr16);
3 alert( arr8[0] ); // 1
4 alert( arr8[1] ); // 232, tried to copy 1000, but can't fit 1000 into 8 bit
```

4. For a numeric argument `length` – creates the typed array to contain that many elements. Its byte length will be `length` multiplied by the number of bytes in a single item `TypedArray.BYTES_PER_ELEMENT` :

```
1 let arr = new Uint16Array(4); // create typed array for 4 integers
2 alert( Uint16Array.BYTES_PER_ELEMENT ); // 2 bytes per integer
3 alert( arr.byteLength ); // 8 (size in bytes)
```

5. Without arguments, creates an zero-length typed array.

We can create a `TypedArray` directly, without mentioning `ArrayBuffer` . But a view cannot exist without an underlying `ArrayBuffer` , so gets created automatically in all these cases except the first one (when provided).

To access the `ArrayBuffer` , there are properties:

- `arr.buffer` – references the `ArrayBuffer` .
- `arr.byteLength` – the length of the `ArrayBuffer` .

So, we can always move from one view to another:

```
1 let arr8 = new Uint8Array([0, 1, 2, 3]);
2
3 // another view on the same data
4 let arr16 = new Uint16Array(arr8.buffer);
```

Here's the list of typed arrays:

- `Uint8Array` , `Uint16Array` , `Uint32Array` – for integer numbers of 8, 16 and 32 bits.
  - `Uint8ClampedArray` – for 8-bit integers, “clamps” them on assignment (see below).
- `Int8Array` , `Int16Array` , `Int32Array` – for signed integer numbers (can be negative).
- `Float32Array` , `Float64Array` – for signed floating-point numbers of 32 and 64 bits.



### No `int8` or similar single-valued types

Please note, despite of the names like `Int8Array` , there's no single-value type like `int` , or `int8` in JavaScript.

That's logical, as `Int8Array` is not an array of these individual values, but rather a view on `ArrayBuffer` .

## Out-of-bounds behavior

What if we attempt to write an out-of-bounds value into a typed array? There will be no error. But extra bits are cut-off.

For instance, let's try to put 256 into `Uint8Array` . In binary form, 256 is `100000000` (9 bits), but `Uint8Array` only provides 8 bits per value, that makes the available range from 0 to 255.

For bigger numbers, only the rightmost (less significant) 8 bits are stored, and the rest is cut off:

### 8-bit integer

1 `00000000` 256

So we'll get zero.

For 257, the binary form is `100000001` (9 bits), the rightmost 8 get stored, so we'll have `1` in the array:

### 8-bit integer

1 `00000001` 257

In other words, the number modulo  $2^8$  is saved.

Here's the demo:

```
1 let uint8array = new Uint8Array(16);
2
3 let num = 256;
4 alert(num.toString(2)); // 100000000 (binary representation)
5
6 uint8array[0] = 256;
7 uint8array[1] = 257;
8
9 alert(uint8array[0]); // 0
10 alert(uint8array[1]); // 1
```



`Uint8ClampedArray` is special in this aspect, its behavior is different. It saves 255 for any number that is greater than 255, and 0 for any negative number. That behavior is useful for image processing.

## TypedArray methods

`TypedArray` has regular `Array` methods, with notable exceptions.

We can iterate, `map`, `slice`, `find`, `reduce` etc.

There are few things we can't do though:

- No `splice` – we can't “delete” a value, because typed arrays are views on a buffer, and these are fixed, contiguous areas of memory. All we can do is to assign a zero.
- No `concat` method.

There are two additional methods:

- `arr.set(fromArr, [offset])` copies all elements from `fromArr` to the `arr`, starting at position `offset` (0 by default).
- `arr.subarray([begin, end])` creates a new view of the same type from `begin` to `end` (exclusive). That's similar to `slice` method (that's also supported), but doesn't copy anything – just creates a new view, to operate on the given piece of data.

These methods allow us to copy typed arrays, mix them, create new arrays from existing ones, and so on.

## DataView

`DataView` is a special super-flexible “untyped” view over `ArrayBuffer`. It allows to access the data on any offset in any format.

- For typed arrays, the constructor dictates what the format is. The whole array is supposed to be uniform. The *i*-th number is `arr[i]`.

- With `DataView` we access the data with methods like `.getUint8(i)` or `.getUint16(i)`. We choose the format at method call time instead of the construction time.

The syntax:

```
1 new DataView(buffer, [byteOffset], [byteLength])
```

- **buffer** – the underlying `ArrayBuffer`. Unlike typed arrays, `DataView` doesn't create a buffer on its own. We need to have it ready.
- **byteOffset** – the starting byte position of the view (by default 0).
- **byteLength** – the byte length of the view (by default till the end of `buffer`).

For instance, here we extract numbers in different formats from the same buffer:

```
1 // binary array of 4 bytes, all have the maximal value 255
2 let buffer = new Uint8Array([255, 255, 255, 255]).buffer;
3
4 let dataView = new DataView(buffer);
5
6 // get 8-bit number at offset 0
7 alert( dataView.getUint8(0) ); // 255
8
9 // now get 16-bit number at offset 0, it consists of 2 bytes, together interpr
10 alert( dataView.getUint16(0) ); // 65535 (biggest 16-bit unsigned int)
11
12 // get 32-bit number at offset 0
13 alert( dataView.getUint32(0) ); // 4294967295 (biggest 32-bit unsigned int)
14
15 dataView.setUint32(0, 0); // set 4-byte number to zero, thus setting all byte
```

`DataView` is great when we store mixed-format data in the same buffer. E.g we store a sequence of pairs (16-bit integer, 32-bit float). Then `DataView` allows to access them easily.

## Summary

`ArrayBuffer` is the core object, a reference to the fixed-length contiguous memory area.

To do almost any operation on `ArrayBuffer`, we need a view.

- It can be a `TypedArray` :
  - `Uint8Array`, `Uint16Array`, `Uint32Array` – for unsigned integers of 8, 16, and 32 bits.
  - `Uint8ClampedArray` – for 8-bit integers, “clamps” them on assignment.
  - `Int8Array`, `Int16Array`, `Int32Array` – for signed integer numbers (can be negative).
  - `Float32Array`, `Float64Array` – for signed floating-point numbers of 32 and 64 bits.
- Or a `DataView` – the view that uses methods to specify a format, e.g. `getUint8(offset)`.

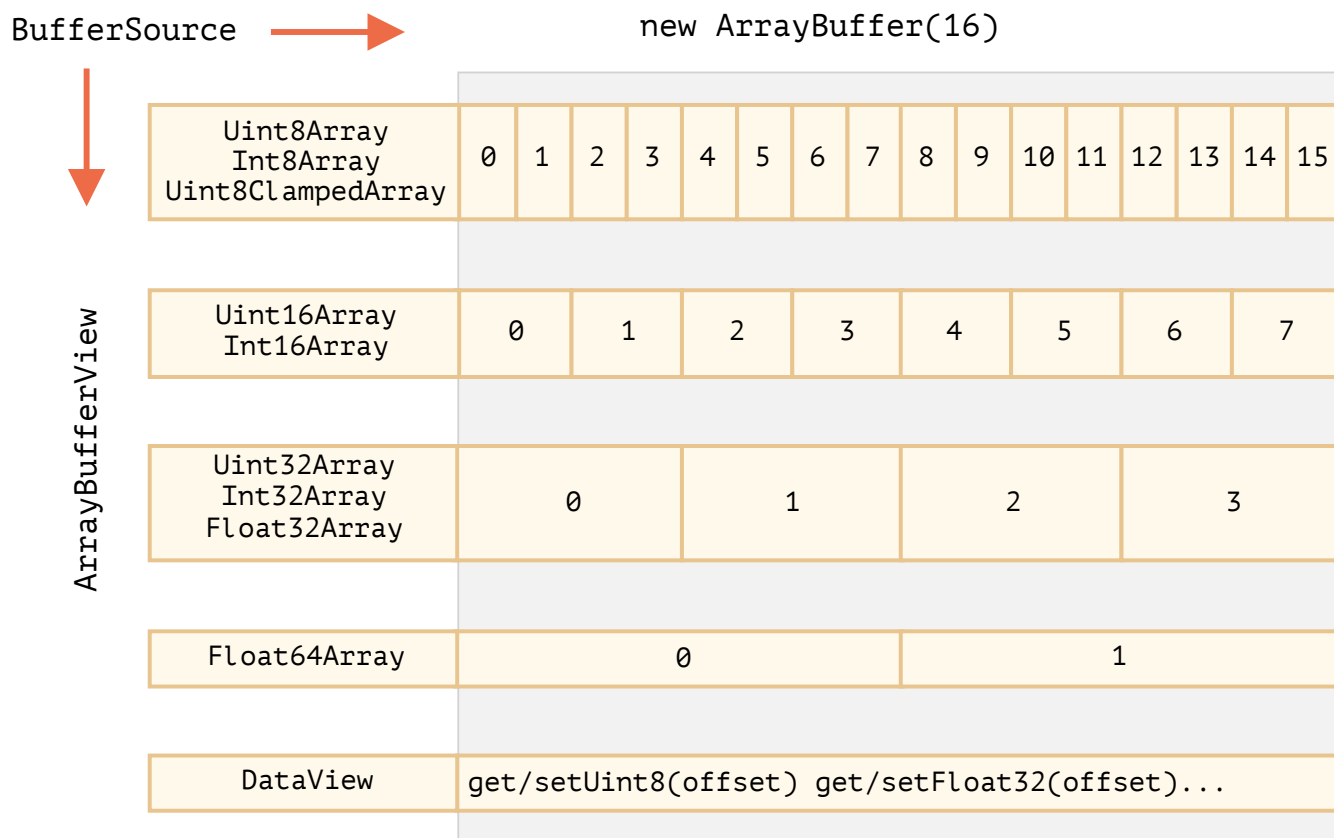
In most cases we create and operate directly on typed arrays, leaving `ArrayBuffer` under cover, as a “common discriminator”. We can access it as `.buffer` and make another view if needed.

There are also two additional terms, that are used in descriptions of methods that operate on binary data:

- `ArrayBufferView` is an umbrella term for all these kinds of views.
- `BufferSource` is an umbrella term for `ArrayBuffer` or `ArrayBufferView`.

We'll see these terms in the next chapters. `BufferSource` is one of the most common terms, as it means "any kind of binary data" – an `ArrayBuffer` or a view over it.

Here's a cheatsheet:



## ✓ Tasks

### Concatenate typed arrays [↗](#)

Given an array of `Uint8Array`, write a function `concat(arrays)` that returns a concatenation of them into a single array.

[Open a sandbox with tests.](#)

[solution](#)



Previous lesson

Next lesson



## Comments

- If you have suggestions what to improve - please [submit a GitHub issue](#) or a pull request instead of commenting.
- If you can't understand something in the article – please elaborate.
- To insert a few words of code, use the `<code>` tag, for several lines – use `<pre>` , for more than 10 lines – use a sandbox ([plnkr](#), [JSBin](#), [codepen](#)...)